

# Systems Programming



## The Stat System Call

# Shell revisited : Redirecting stderr



- 3.3 Saving stderr to a file
- `gcc prog.c 2> compile-errors.txt`
- 3.5 Saving stderr to stdout
- `gcc prog.c 2> &1`
- 3.5 Saving stderr and stdout to a file
- `rm -f $(find / -name core) &> /dev/null`

<http://tldp.org/HOWTO/Bash-Prog-Intro-HOWTO-3.html>

# /proc (Proc File System)



- “an interface to kernel data structures in a form that looks like files and directories”
- /proc/<pid> - directory of info about running process
  - /proc/1 – information on init
- /proc/cpuinfo
- /proc/sys/fs/file-max



```
DIR *dirp;  
struct dirent *dp;  
Boolean isCurrent; /* True if 'dirpath' is "." */  
isCurrent = strcmp(dirpath, ".") == 0;  
dirp = opendir(dirpath);  
if (dirp == NULL) {  
    errMsg("opendir failed on '%s'", dirpath);  
    return;  
}
```



```
/* For each entry in this directory, print directory +
filename */
for (;;) {
    errno = 0; // To distinguish error from end-of-directory
    dp = readdir(dirp);
    if (dp == NULL)
        break;

    if (strcmp(dp->d_name, ".") == 0 || strcmp(dp->d_name,
"..") == 0)
        continue;                /* Skip . and .. */

    if (!isCurrent)
        printf("%s/", dirpath);
    printf("%s\n", dp->d_name);
}
```



```
if (errno != 0)
    errExit("readdir");
if (closedir(dirp) == -1)
    errMsg("closedir");
```

# TLPI Source Code for practice



**Please Download the book's source code:**

<https://github.com/posborne/linux-programming-interface-exercises>

or

<https://projects.ui.ac.id/attachments/download/7181/linux-programming-interface-exercises-master.zip>

**Contents in “tlpi-dist” directory:**

```
root> ls
```

```
acl                                getopt                            pmsg                             svmsg
...
daemons  Makefile.inc.MacOSX        pshm                             tty README
dirs_links  Makefile.inc.Solaris          pty                             users_groups
...
files                                pgsjc                           sockets
filesys                             pipes                            svipc
```

**NB: Please Compile the source codes using “make”**

# TLPI/dirs\_links



- `cd dir_links`

- `ls`

`bad_symlink.c`

`list_files_readdir_r.c`

`t_dirbasename.c`

`file_type_stats.c`

`Makefile`

`t_unlink.c`

`list_files.c`

`nftw_dir_tree.c`

`view_symlink.c`



# list\_files\_readdir\_r.c : listFiles()



```
DIR *dirp;
Boolean isCurrent; /* True if 'dirpath'
is "." */
struct dirent *result, *entryp;
int nameMax;

isCurrent = strcmp(dirpath, ".") == 0;
nameMax = pathconf(dirpath,
_PC_NAME_MAX);
if (nameMax == -1) /* Indeterminate or
error */
    nameMax = 255; /* So take a guess */
entryp = malloc(offsetof(struct dirent,
d_name) + nameMax + 1);
if (entryp == NULL)
    errExit("malloc");
dirp = opendir(dirpath);
if (dirp == NULL) {
    errMsg("opendir failed on '%s'",
dirpath);
    return;
}
```

```
for (;;) {
    errno = readdir_r(dirp, entryp,
&result);
    if (errno != 0)
        errExit("readdir_r");

    if (result == NULL) /* End of stream */
        break;

    /* Skip . and .. */

    if (strcmp(entryp->d_name, ".") == 0 ||
        strcmp(entryp->d_name, "..")
== 0)
        continue;

    /* Print directory + filename */

    if (!isCurrent) printf("%s/", dirpath);
    printf("%s\n", entryp->d_name);
}

if (closedir(dirp) == -1)
    errMsg("closedir");
```

## Example 18-2. Scanning a directory



`readdir_r` - reentrant version of `readdir`

The `list_files.c` use `readdir`

The `list_files_readdir_r.c` use `readdir_r`

Want to find out `readdir` vs `readdir_r` ?

Please call the man

`# man readdir`

# So where does ls get all that information on files?



The inode table contains information on files: permissions, times, etc.

To allow programmers to get at this information Unix provides the stat system call

# man 2 stat



STAT(2)      Linux Programmer's Manual

STAT(2)

## NAME

stat, fstat, lstat - get file status

## SYNOPSIS

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <unistd.h>
```

```
int stat(const char *path, struct stat *buf); /* file path*/
```

```
int fstat(int fd, struct stat *buf); /* file descriptors */
```

```
int lstat(const char *path, struct stat *buf); /* symbolic link path */
```

# Stat structure (man 2 stat)



```
struct stat {  
    dev_t    st_dev;        /* ID of device containing file */  
    ino_t    st_ino;        /* inode number */  
    mode_t   st_mode;       /* file type and permissions */  
    nlink_t  st_nlink;      /* number of hard links */  
    uid_t    st_uid;        /* user ID of owner */  
    gid_t    st_gid;        /* group ID of owner */  
    dev_t    st_rdev;       /* device ID (if special file) */  
    off_t    st_size;       /* total size, in bytes */  
    blksize_t st_blksize;   /* blocksize for file system I/O */  
    blkcnt_t st_blocks;     /* number of 512B blocks allocated */  
    time_t   st_atime;      /* time of last access */  
    time_t   st_mtime;      /* time of last modification */  
    time_t   st_ctime;      /* time of last status change */  
};
```

# st\_mode macros (man 2 stat)



**S\_ISREG(m)** is it a regular file?

**S\_ISDIR(m)** directory?

**S\_ISCHR(m)** character device?

- A Character ('c') Device is one with which the Driver communicates by sending and receiving single characters (bytes, octets).
- Examples for Character Devices: serial ports, parallel ports, sounds cards.

**S\_ISBLK(m)** block device?

- A Block ('b') Device is one with which the Driver communicates by sending entire blocks of data.
- Examples for Block Devices: hard disks,

**S\_ISFIFO(m)** FIFO (named pipe)?

**S\_ISLNK(m)** symbolic link? (Not in POSIX.1-1996.)

**S\_ISSOCK(m)** socket? (Not in POSIX.1-1996.)

# st\_mode Flags (man 2 stat)



- `ls /usr/include/sys`

...

- `less /usr/include/sys/stat.h`

...

```
#define S_IRUSR  __S_IREAD    /* Read by owner.  */
#define S_IWUSR  __S_IWRITE   /* Write by owner. */
#define S_IXUSR  __S_IEXEC    /* Execute by owner. */
```

...

```
#define S_IRGRP (S_IRUSR >> 3) /* Read by group.  */
```

# Finding Foo



```
root> man -s 5 ar
```

```
No manual entry for ar
```

```
See 'man 7 undocumented'
```

- Finding the command foo try the following options :

foo --help, foo -h, foo -?

info foo

whatis foo, apropos foo

dpkg --listfiles foo, dpkg --search foo

locate '\*foo\*'

find / -name '\*foo\*'

check /usr/share/doc/foo, /usr/lib/foo.



# Chapter 4. File I/O: The Universal I/O Model



- Unix I/O = files
- File descriptors - index into open file table
- `_iob[_NFILES]` struct - per process open file table
- standard descriptors
  - 0 = `STDIN_FILENO`
  - 1 = `STDOUT_FILENO`
  - 2 = `STDERR_FILENO`
- Remapping - duplicate [Read Ch 5.4]

# Four Key System Calls Performing file I/O



- `fd = open (pathname, flags, mode)`
- `numread = read (fd, buffer, count)`
- `numwritten = write (fd, buffer, count)`
- `status = close (fd)`

# Unix I/O system calls



- `ssize_t read(int fildes, void *buf, size_t nbyte);`
- `ssize_t write(int fildes, const void *buf, size_t nbyte);`
- `int open(const char *path, int oflag, /* mode_t mode */...);`
- `int close(int fildes);`
- `int creat(const char *path, mode_t mode);`
  - equivalent to `open(path, O_WRONLY | O_CREAT | O_TRUNC, mode)`
- `off_t lseek(int fildes, off_t offset, int whence);`

# TLPI/fileio/copy.c



```
...
int
main(int argc, char *argv[])
{
    int inputFd, outputFd, openFlags;
    mode_t filePerms;
    ssize_t numRead;
    char buf[BUF_SIZE];

    if (argc != 3 || strcmp(argv[1],
        "--help") == 0)
        usageErr("%s old-file new-file\n",
            argv[0]);

    /* Open input and output files */
    inputFd = open(argv[1], O_RDONLY);
    if (inputFd == -1)
        errExit("opening file %s",
            argv[1]);
```

```
    openFlags = O_CREAT | O_WRONLY | O_TRUNC;
    filePerms = S_IRUSR | S_IWUSR | S_IRGRP |
        S_IWGRP | S_IROTH | S_IWOTH; /* rw-rw-rw- */
    outputFd = open(argv[2], openFlags,
        filePerms);
    if (outputFd == -1)
        errExit("opening file %s", argv[2]);

    /* Transfer data until we encountered EOF
    or an error */
    while ((numRead = read(inputFd, buf,
        BUF_SIZE)) > 0)
        if (write(outputFd, buf, numRead) !=
            numRead)
            fatal("couldn't write whole
                buffer");

    if (numRead == -1)
        errExit("read");
    if (close(inputFd) == -1)
        errExit("close input");
    if (close(outputFd) == -1)
        errExit("close output");
    exit(EXIT_SUCCESS);
}
```

# Universality of I/O



`$ ./copy test test.old` -- Copy a regular file

`$ ./copy a.txt /dev/tty`

`$ ./copy /dev/tty b.txt`

`$ ./copy /dev/pts/16 /dev/tty`

`/dev/null`

# OPEN



## NAME

open, creat - open and possibly create a file or device

## SYNOPSIS

```
#include <sys/types.h>
```

```
#include <sys/stat.h>
```

```
#include <fcntl.h>
```

```
int open(const char *pathname, int flags);
```

```
int open(const char *pathname, int flags, mode_t mode);
```

```
int creat(const char *pathname, mode_t mode);
```



*Table 4-2. File access modes*

Access mode	Description
O_RDONLY	Open the file for reading only
O_WRONLY	Open the file for writing only
O_RDWR	Open the file for both reading and writing

# Table 4.3 Flags for open



Flag	Purpose	SUS?
O_RDONLY	Open for reading only	v3
O_WRONLY	Open for writing only	v3
O_RDWR	Open for reading and writing	v3
O_CLOEXEC	Set the close-on-exec flag	v4
O_CREAT	Create file if it doesn't already exist	v3
O_DIRECT	File I/ O bypasses buffer cache	
O_DIRECTORY	Fail if pathname is not a directory	v4
O_EXCL With O_CREAT	create file exclusively	v3



# Btw, what is SUS ?



## Is it a Sus Cake (Kue soes merdeka) ? :)

### Chapter 3.6 - Portability Issue



# Table 4.3 Flags for open



Flag	Purpose	SUS?
O_RDONLY	Open for reading only	v3
O_WRONLY	Open for writing only	v3
O_RDWR	Open for reading and writing	v3
O_CLOEXEC	Set the close-on-exec flag (since Linux 2.6.23)	v4
O_CREAT	Create file if it doesn't already exist	v3
O_DIRECT	File I/O bypasses buffer cache	
O_DIRECTORY	Fail if <i>pathname</i> is not a directory	v4
O_EXCL	With O_CREAT: create file exclusively	v3
O_LARGEFILE	Used on 32-bit systems to open large files	
O_NOATIME	Don't update file last access time on <i>read()</i> (since Linux 2.6.8)	
O_NOCTTY	Don't let <i>pathname</i> become the controlling terminal	v3
O_NOFOLLOW	Don't dereference symbolic links	v4
O_TRUNC	Truncate existing file to zero length	v3
O_APPEND	Writes are always appended to end of file	v3
O_ASYNC	Generate a signal when I/O is possible	
O_DSYNC	Provide synchronized I/O data integrity (since Linux 2.6.33)	v3
O_NONBLOCK	Open in nonblocking mode	v3
O_SYNC	Make file writes synchronous	v3

Kerrisk, Michael  
(2011-02-11). *The Linux  
Programming Interface: A  
Linux and UNIX System  
Programming Handbook.*

# Listing 4.2 - open file for Reading



Please find the code at:  
[tlpi-dist/fileio/copy.c](http://tlpi-dist/fileio/copy.c)

```
/* Open input and output files */  
inputFd = open(argv[1], O_RDONLY);  
if (inputFd == -1)  
    errExit("opening file %s", argv[1]);  
  
openFlags = O_CREAT | O_WRONLY | O_TRUNC;  
filePerms = S_IRUSR | S_IWUSR | S_IRGRP | S_IWGRP |  
            S_IROTH | S_IWOTH;          /* rw-rw-rw- */  
outputFd = open(argv[2], openFlags, filePerms);  
if (outputFd == -1)  
    errExit("opening file %s", argv[2]);
```

File Descriptor  
to read source file

File  
Descriptor  
to write  
destination  
file

# Errors from open()



- Retval = -1
- Errno
  - EACCES
  - EISDIR
  - EMFILE
  - ENFILE
  - ENOENT – file does not exist and O\_CREAT not specified
  - EROFS – file system read only
  - ETXTBSY – executable that is currently running

# *Creat()* system call



- **Creat equivalent to open with certain flags**
- **`fd = open( pathname, O_WRONLY | O_CREAT | O_TRUNC, mode);`**
- **Less powerful compared to `open()`, thus obsolete**

# The Read system call



## NAME

read - read from a file descriptor

## SYNOPSIS

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t count);
```

## DESCRIPTION

read() attempts to read up to count bytes from file descriptor fd into the buffer starting at buf.

If count is zero, read() returns zero and has no other results.

If count is greater than SSIZE\_MAX, the result is unspecified.

# Write, Close



```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t count);
```

DESCRIPTION - write() writes up to count bytes from the buffer pointed buf to the file referred to by the file descriptor fd.

SYNOPSIS

.....

```
int close(int fd);
```

DESCRIPTION - close() closes a file descriptor, so that it no longer refers to any file and may be reused.

# Changing the file OFFSET - lseek



## NAME

**lseek** - reposition read/write file offset

## SYNOPSIS

```
#include <sys/types.h>
```

```
#include <unistd.h>
```

```
off_t lseek(int fd, off_t offset, int whence);
```

**DESCRIPTION** - The **lseek()** function repositions the offset of the open file associated with the file descriptor **fd** to the argument **offset** according to the directive **whence** as follows:

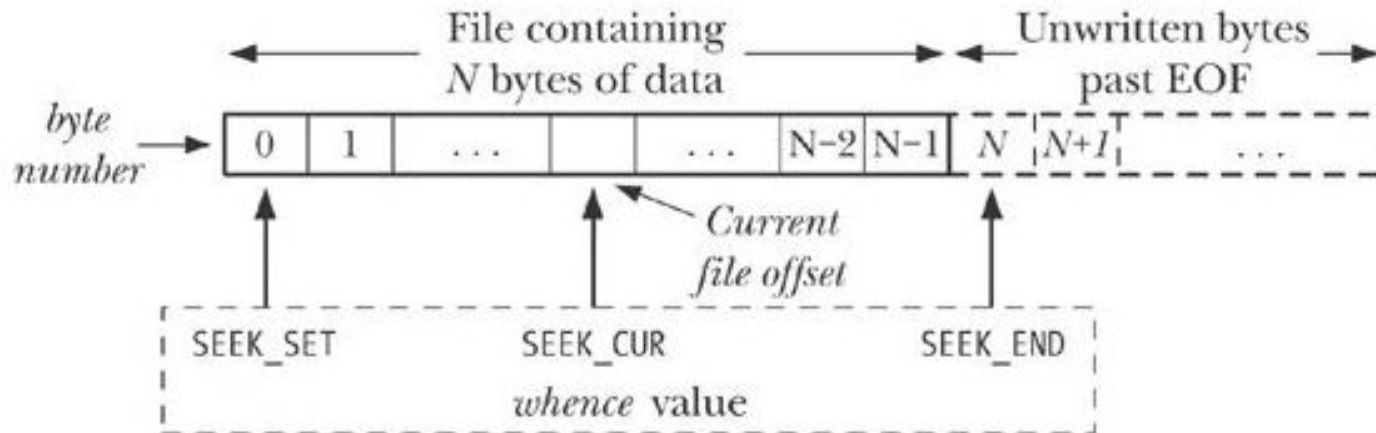
**SEEK\_SET** - The offset is set to **offset** bytes.

**SEEK\_CUR** - The offset is set to its current location plus **offset** bytes.

**SEEK\_END** - The offset is set to the size of the file plus **offset** bytes.



# Figure 4.1 illustrating lseek



*Figure 4-1. Interpreting the whence argument of lseek()*

```
lseek(fd, 0, SEEK_SET); /* Start of file */
lseek(fd, 0, SEEK_END); /* Next byte after the EOF */
lseek(fd, -1, SEEK_END); /* Last byte of file */
lseek(fd, -10, SEEK_CUR); /* 10 bytes prior to current */
lseek(fd, 100, SEEK_CUR); /* 100 bytes past of current */
```



IOCTL(2)

Linux Programmer's Manual

IOCTL(2)

NAME

ioctl - control device

SYNOPSIS

```
#include <sys/ioctl.h>
```

```
int ioctl(int d, int request, ...);
```

**DESCRIPTION** - The `ioctl()` function manipulates the underlying device parameters of special files.

# Example IOCTL usage



- Example 1: IOCTL on CDROM drive

First, create the file descriptor

```
fd = open('/dev/cdrom', O_RDONLY |  
O_NONBLOCK);
```

Eject:

```
ioctl(fd, CDROMEJECT, 0);
```

Close Tray:

```
ioctl(fd, CDROMCLOSETRAY, 0);
```



# QA

# Example IOCTL usage



- Example 2: Keyboard LED

```
#include <stdio.h>
#include <linux/kd.h>

int main(int argc, char *argv[])
{
    ioctl(1, KDSETLED, NUM_LED);
    return 0;
}
```

Where is the open file descriptor ?

Why accessing the Keyboard can use FD = 1 ?

# More flags



- **O\_EXCL** With **O\_CREAT**: create file exclusively v3
- **O\_LARGEFILE** Used on 32-bit systems to open large files
- **O\_NOATIME** Don't update file last access time on read() (since Linux 2.6.8)
- **O\_NOCTTY** Don't let pathname become the controlling terminal v3
- **O\_NOFOLLOW** Don't dereference symbolic links v4
- **O\_TRUNC** Truncate existing file to zero length v3
- **O\_APPEND** Writes are always appended to end of file v3
- **O\_ASYNC** Generate a signal when I/ O is possible
- **O\_DSYNC** Provide synchronized I/ O data integrity (since Linux 2.6.33) v3
- **O\_NONBLOCK** Open in nonblocking mode v3
- **O\_SYNC** Make file writes synchronous v3
-

# Open file structures



- per process table
- open file table
- vtable